

FIXED CONTEXT

Max 512 tokens
Long context lost

ENCODER-ONLY

No autoregressive gen
Not for text creation

HIGH COST

Heavy compute
110M–340M params

WEAK REASONING

Pattern-based only
Poor multi-step logic

DATA BIAS

 Learns & amplifies
dataset bias

TASK-SPECIFIC

Needs fine-tuning
Not plug-and-play